

HIGH PERFORMANCE COMPUTING: TOWARDS BETTER PERFORMANCE PREDICTIONS AND EXPERIMENTS

Tom Cornebize


2 June 2021, PhD defense



NO SCIENCE WITHOUT COMPUTING



Arithmomètre (1851)



ENIAC (1945)

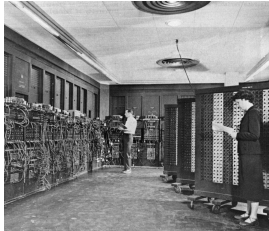


Fugaku (2021)

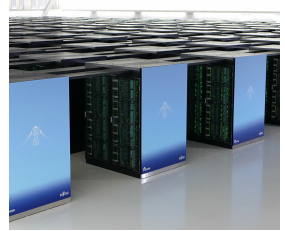
NO SCIENCE WITHOUT COMPUTING



Arithmomètre (1851)



ENIAC (1945)

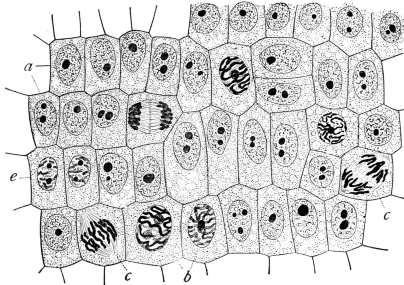


Fugaku (2021)

Last decades:

- Exponential **performance** improvements (e.g. sequencing an entire human genome costed \$100,000,000 in 2001, \$1000 now)
- At the price of **complexity** (both software and hardware)

EXPERIMENTAL STUDY OF COMPUTER PERFORMANCE



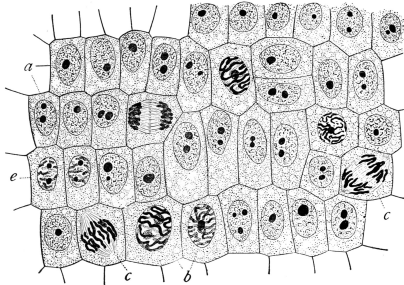
Similar to natural sciences

Complexity \Rightarrow Variability and Opacity

\Rightarrow No perfect model

\Rightarrow Need for **experiments**

EXPERIMENTAL STUDY OF COMPUTER PERFORMANCE



Similar to natural sciences

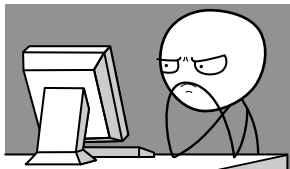
Complexity \Rightarrow Variability and Opacity

\Rightarrow No perfect model

\Rightarrow Need for **experiments**

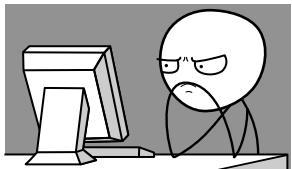
Experiments can be carried in **reality** or in **simulation**

Typical Performance Evaluation Questions (Given my application and a supercomputer)



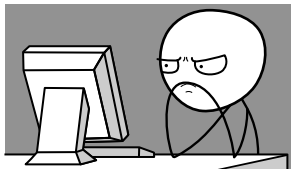
- **Before** running
 - How many nodes?
 - For how long?
 - Which parameters?

Typical Performance Evaluation Questions (Given my application and a supercomputer)



- **Before** running
 - How many nodes?
 - For how long?
 - Which parameters?
- **After** running
 - Performance as “expected”?
 - Problem in the app or the platform?


Typical Performance Evaluation Questions (Given my application and a supercomputer)



- **Before** running
 - How many nodes?
 - For how long?
 - Which parameters?
- **After** running
 - Performance as “expected”?
 - Problem in the app or the platform?

So many large-scale runs, solely to tune performance?!?

Typical Performance Evaluation Questions (Given my application and a supercomputer)



- **Before** running
 - How many nodes?
 - For how long?
 - Which parameters?
- **After** running
 - Performance as “expected”?
 - Problem in the app or the platform?

So many large-scale runs, solely to tune performance?!?

Holy Grail: Predictive Simulation on a “Laptop”

Capture the **whole application** and **platform complexity**

Initial goal: **predict** the performance of
a parallel application

Initial goal: **predict** the performance of a parallel application

Thesis contributions (towards this goal)

- Case study: High Performance Linpack (HPL)
- Extensive (in)validation, comparing simulations with reality
- Modeling correctly the platform variability is key

Initial goal: **predict** the performance of a parallel application

Thesis contributions (towards this goal)

- Case study: High Performance Linpack (HPL)
- Extensive (in)validation, comparing simulations with reality
- Modeling correctly the platform variability is key

Thesis contributions (made on the way)

- Automation (of experiments, statistical analyzes, etc.)
- Experiment methodology, to bias or not to bias
- Performance tests, to detect eventual platform changes

Initial goal: **predict** the performance of a parallel application

Thesis contributions (towards this goal)

- Case study: High Performance Linpack (HPL)
- Extensive (in)validation, comparing simulations with reality
- Modeling correctly the platform variability is key


Thesis contributions (made on the way)

- Automation (of experiments, statistical analyzes, etc.)
- Experiment methodology, to bias or not to bias
- Performance tests, to detect eventual platform changes

PERFORMANCE PREDICTION THROUGH SIMULATION

SIM(EM)ULATION: THE SMPI APPROACH



Full reimplementation of MPI on top of  SIMGRID

- C/C++/F77/F90 codes run **unmodified out of the box**
- Simply replace mpicc/mpirun by smpicc/smpirun



SIM(EM)ULATION: THE SMPI APPROACH



Full reimplementation of MPI on top of 

- C/C++/F77/F90 codes run **unmodified out of the box**
- Simply replace mpicc/mpirun by smpicc/smpirun




Emulation: how?

- Computations run for real on a laptop
- Communications are faked, good fluid network models
- **Performance model** for the target platform

SIM(EM)ULATION: THE SMPI APPROACH



Full reimplementation of MPI on top of  SIMGRID

- C/C++/F77/F90 codes run **unmodified out of the box**
- Simply replace mpicc/mpirun by smpicc/smpirun




Emulation: how?

- Computations run for real on a laptop
- Communications are faked, good fluid network models
- **Performance model** for the target platform

Contribution: Skip the expensive computations (mostly **dgemm**) and replace them by performance models

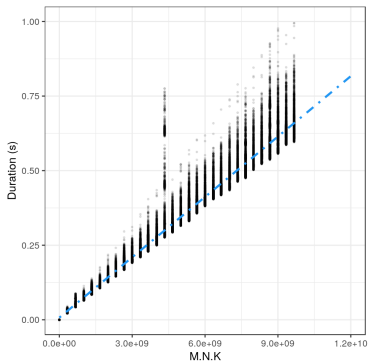
MODELING COMPUTATIONS

$\text{dgemm}(M, N, K) =$



MODELING COMPUTATIONS


$$\text{dgemm}(M, N, K) = \alpha.M.N.K$$



MODELING COMPUTATIONS

$$\text{dgemm}_i(M, N, K) = \underbrace{\alpha_i \cdot M \cdot N \cdot K}_{\text{per host}}$$

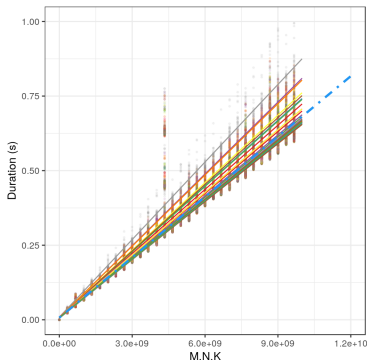
Different color \Rightarrow different host



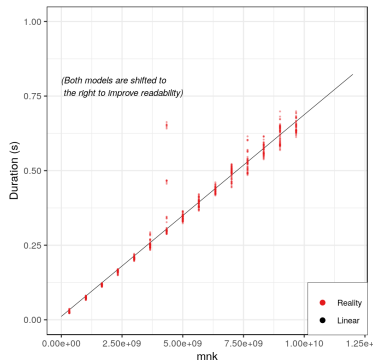
MODELING COMPUTATIONS

$$\text{dgemm}_i(M, N, K) = \underbrace{\alpha_i \cdot M \cdot N \cdot K}_{\text{per host}}$$

Different color \Rightarrow different host




For a particular host




MODELING COMPUTATIONS

$$\text{dgemm}_i(M, N, K) = \underbrace{\alpha_i \cdot M \cdot N \cdot K}_{\text{per host}} + \underbrace{\beta_i \cdot M \cdot N + \gamma_i \cdot N \cdot K + \dots}_{\text{polynomial model}}$$

Different color \Rightarrow different host



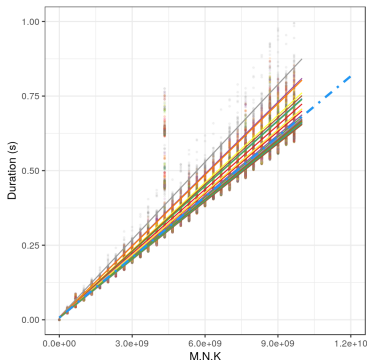
For a particular host



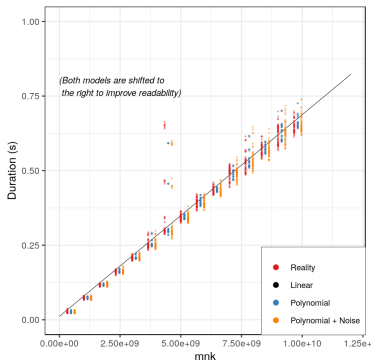
MODELING COMPUTATIONS

$$\text{dgemm}_i(M, N, K) = \underbrace{\alpha_i \cdot M \cdot N \cdot K}_{\text{per host}} + \underbrace{\beta_i \cdot M \cdot N + \gamma_i \cdot N \cdot K + \dots}_{\text{polynomial model}} + \underbrace{\mathcal{N}(0, \alpha'_i \cdot M \cdot N \cdot K + \dots)}_{\text{polynomial noise}}$$

Different color \Rightarrow different host



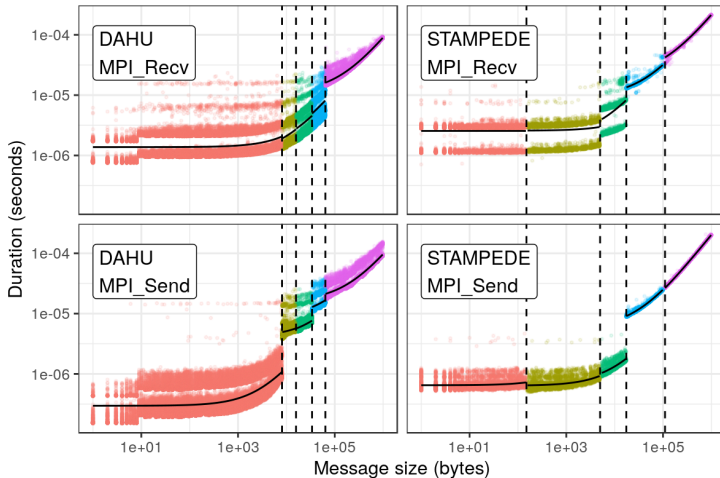
For a particular host



Hand-crafted non-blocking collective operations intertwined with computations

MODELING COMMUNICATIONS

Hand-crafted non-blocking collective operations intertwined with computations



Experimental biases when measuring `dgemm` or MPI durations
Effect on durations, but also other metrics (e.g. CPU frequency)

Experimental biases when measuring `dgemm` or MPI durations

Effect on durations, but also other metrics (e.g. CPU frequency)

- Sampling method for generating the sequence of calls

Experimental biases when measuring `dgemm` or MPI durations

Effect on durations, but also other metrics (e.g. CPU frequency)

- Sampling method for generating the sequence of calls
- Sample itself (for a given sampling method)

Experimental biases when measuring `dgemm` or MPI durations

Effect on durations, but also other metrics (e.g. CPU frequency)

- Sampling method for generating the sequence of calls
- Sample itself (for a given sampling method)
- Interferences between computations and communications

Experimental biases when measuring `dgemm` or MPI durations

Effect on durations, but also other metrics (e.g. CPU frequency)

- Sampling method for generating the sequence of calls
- Sample itself (for a given sampling method)
- Interferences between computations and communications
- Content of the matrices used by `dgemm`

Experimental biases when measuring `dgemm` or MPI durations

Effect on durations, but also other metrics (e.g. CPU frequency)


- Sampling method for generating the sequence of calls
- Sample itself (for a given sampling method)
- Interferences between computations and communications
- Content of the matrices used by `dgemm`

Bias may be desirable in some situations

VALIDATING THE PREDICTIONS


INTERNAL BEHAVIOR OF THE APPLICATION

256 MPI ranks, interrupted after the 5th iteration



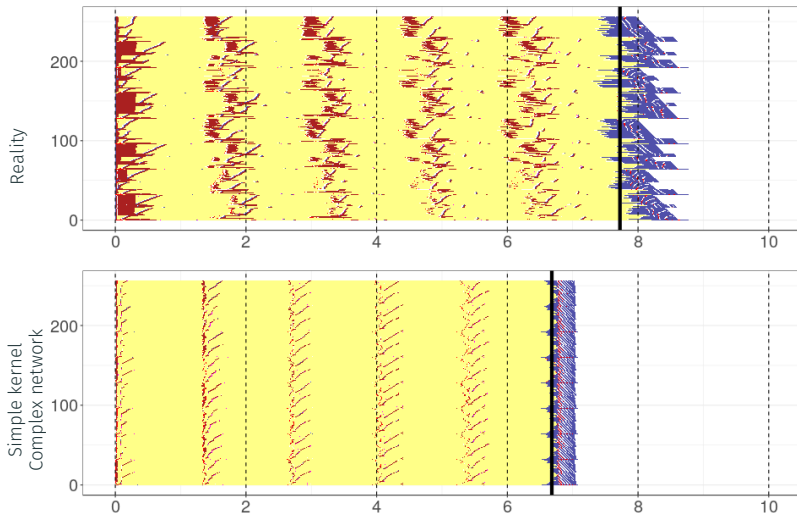
INTERNAL BEHAVIOR OF THE APPLICATION

256 MPI ranks, interrupted after the 5th iteration



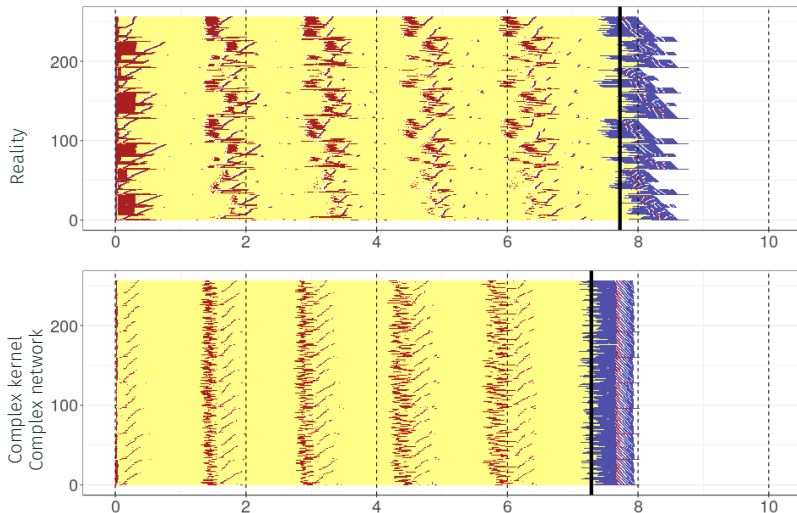
INTERNAL BEHAVIOR OF THE APPLICATION

256 MPI ranks, interrupted after the 5th iteration



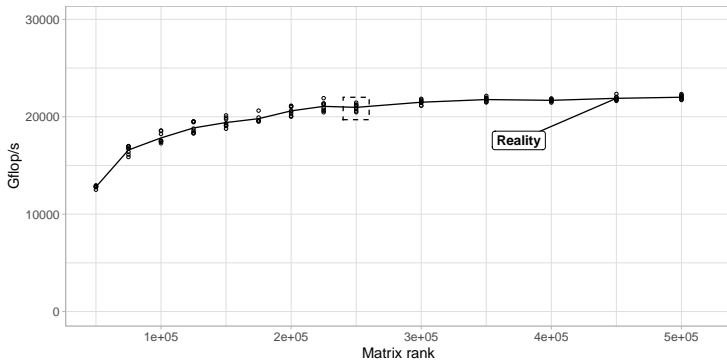
INTERNAL BEHAVIOR OF THE APPLICATION

256 MPI ranks, interrupted after the 5th iteration



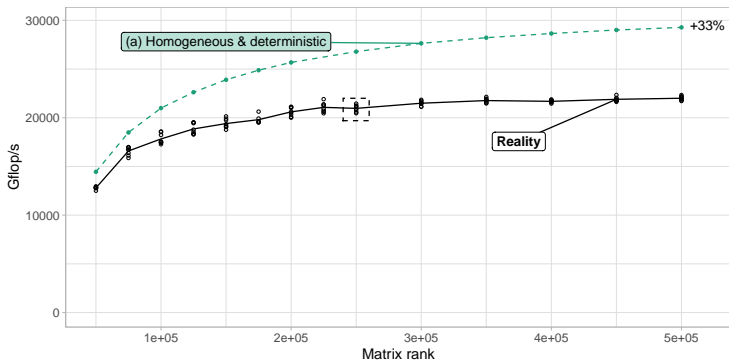
INFLUENCE OF THE PROBLEM SIZE

Now the complete run, with 1024 MPI ranks




INFLUENCE OF THE PROBLEM SIZE

Now the complete run, with 1024 MPI ranks



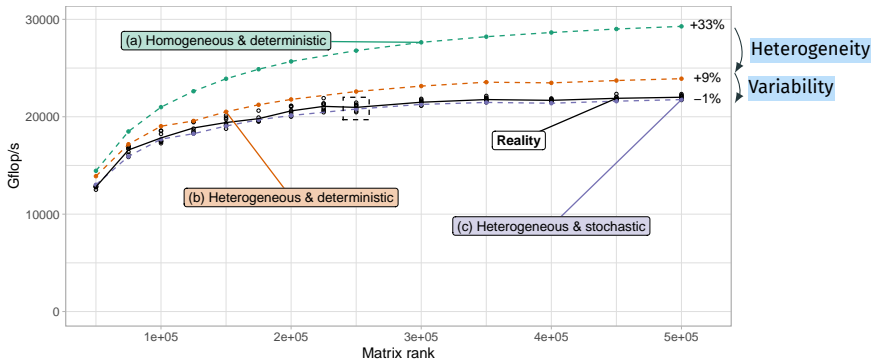
INFLUENCE OF THE PROBLEM SIZE

Now the complete run, with 1024 MPI ranks



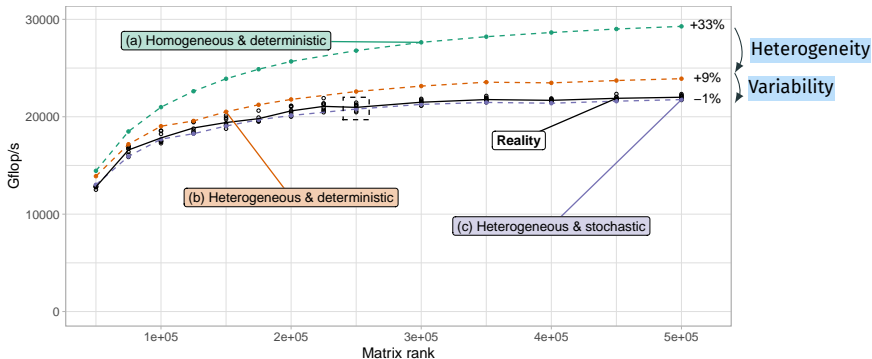
INFLUENCE OF THE PROBLEM SIZE

Now the complete run, with 1024 MPI ranks



INFLUENCE OF THE PROBLEM SIZE

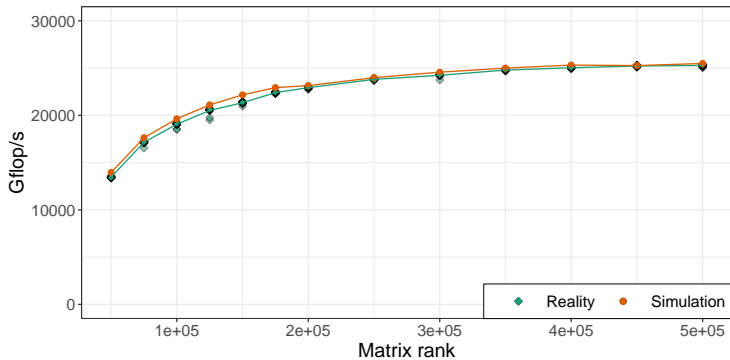
Now the complete run, with 1024 MPI ranks



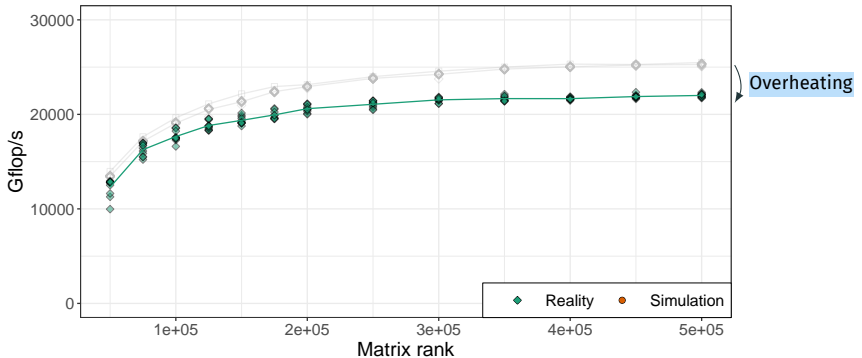
Take-Away Message: accurate prediction

Modeling both **spatial** and **temporal** computation variability is essential

INFLUENCE OF A PLATFORM CHANGE

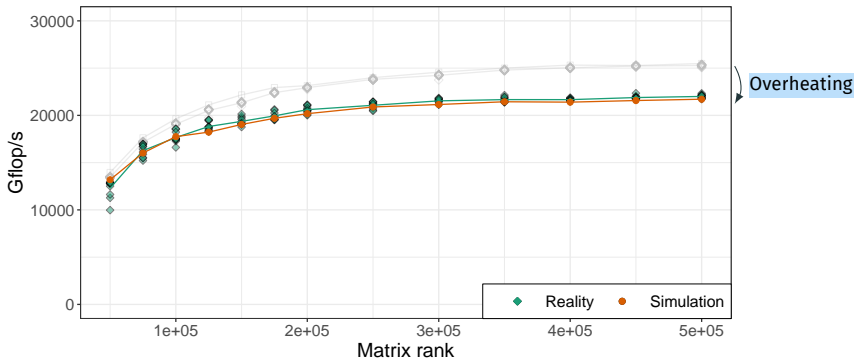


INFLUENCE OF A PLATFORM CHANGE




On four nodes, the cooling system malfunctionned for several weeks

INFLUENCE OF A PLATFORM CHANGE



On four nodes, the cooling system malfunctionned for several weeks

INFLUENCE OF A PLATFORM CHANGE




On four nodes, the cooling system malfunctioned for several weeks

Take-Away Message: Re-measuring `dgemm` durations to generate a new model was enough to account for the platform change

PERFORMANCE TESTS


On a near-daily basis, run the `dgemm` calibration code on
454 nodes (792 CPU) from 12 clusters



On a near-daily basis, run the **dgemm** calibration code on  **Grid'5000**
454 nodes (792 CPU) from 12 clusters

For each CPU, collect:

- average **dgemm** performance
- **dgemm** coefficients of regression (i.e. the model for simulation)

On a near-daily basis, run the **dgemm** calibration code on  **Grid'5000**
454 nodes (792 CPU) from 12 clusters

For each CPU, collect:

- average **dgemm** performance
- **dgemm** coefficients of regression (i.e. the model for simulation)
- average CPU frequency
- average CPU power consumption
- average DRAM power consumption
- average temperature




On a near-daily basis, run the **dgemm** calibration code on
454 nodes (792 CPU) from 12 clusters

For each CPU, collect:

- average **dgemm** performance
- **dgemm** coefficients of regression (i.e. the model for simulation)
- average CPU frequency
- average CPU power consumption
- average DRAM power consumption
- average temperature

Each parameter is **normally distributed** (thanks to CLT)

Given a sequence of **old observations** x_1, \dots, x_n and a **new observation** x_{n+1} , how likely was it to observe x_{n+1} ?



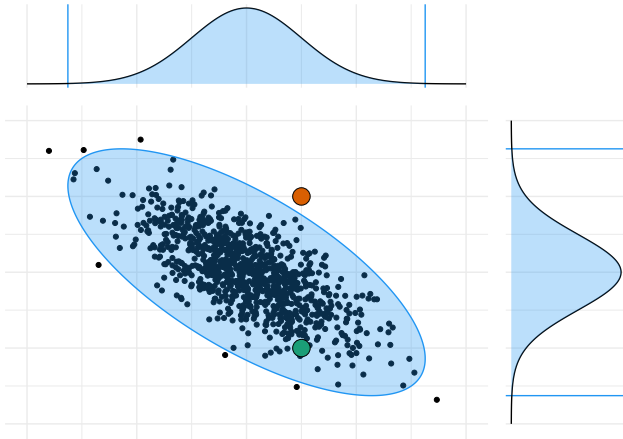
Take the sample mean μ and standard deviation σ of the old observations

$$\mathbb{P}(x_{n+1} \in [\mu - 2\sigma; \mu + 2\sigma]) \approx 95\%$$

FLUCTUATION INTERVAL FOR SEVERAL VARIABLES

With several variables, using their [covariance matrix](#)

Example in dimension 2, with $\mathbb{P}(x_{n+1} \in \text{interval}) \approx 99.5\%$



FLUCTUATION INTERVAL FOR SEVERAL MEASURES

With several measures, using their [average](#) and shrinking the interval

Example with 5 measures (averages represented by crosses)

